# RIYA JANGID

Data Scientist • ML Engineer • Deep Learning Specialist • MLOps

riyajangid02@gmail.com • LinkedIn • GitHub • Jaipur, India

## TECHNICAL SKILLS

**Deep Learning & Neural Networks:** TensorFlow, Keras, CNN, RNN, LSTM, GRU, Word Embeddings, Sequence Modeling, Text Generation, Transfer Learning

**Machine Learning:** NLP (TF-IDF, Word2Vec, N-grams, Tokenization), Random Forest, XGBoost, Logistic Regression, K-Means, SMOTE, Feature Engineering, Hyperparameter Tuning

**MLOps & Deployment:** MLflow, CI/CD Pipelines, Docker, FastAPI, Streamlit, AWS EC2, Model Versioning

**Data & Analytics:** Python, MySQL, Pandas, NumPy, Scikit-learn, EDA, Data Preprocessing, Statistical Analysis, Collaborative Filtering, Cosine Similarity

## EXPERIENCE

**Data Scientist Intern** | Regex Software Services

*Sep 2025 – Present*

- Deployed 5+ production ML systems on AWS EC2 with end-to-end CI/CD pipelines, achieving 94% ROC-AUC in fraud detection and 84% CV score in customer churn prediction across imbalanced real-world datasets.
- Engineered deep learning models including CNN image classifier (92% accuracy), NLP sentiment analyzer (88% accuracy), and LSTM-based sequence models using TensorFlow/Keras; implemented data augmentation and advanced NLP pipelines (TF-IDF, Word2Vec, embeddings) to reduce overfitting by 15%.
- Containerized ML microservices with Docker + FastAPI, cutting deployment time by 60% and enabling real-time inference endpoints with sub-100ms latency under production load.

## PROJECTS

**Next Word Generator with Stacked LSTM** | TensorFlow · Keras · LSTM · GRU · Word Embeddings · NLP

- Built multi-layer stacked LSTM architecture with embedding layer for context-aware text generation; trained on 50K+ sequences using custom tokenization pipeline and optimized embedding dimensions.
- Implemented vocabulary of 10K+ words with optimized 128D embeddings and applied dropout regularization to prevent overfitting; achieved faster convergence through hyperparameter tuning.

**Fraud Detection System** | Random Forest · SMOTE · MLflow · Docker · AWS EC2 · FastAPI ↗ GitHub

- Real-time fraud detection API achieving 94% ROC-AUC on severely imbalanced data (1:99 ratio) using SMOTE oversampling and Random Forest; end-to-end pipeline with MLflow experiment tracking, MySQL audit logging, and Dockerized deployment on AWS EC2.
- Automated CI/CD pipeline eliminates manual deployment steps, enabling zero-downtime model updates with full version traceability and A/B testing capability.

**CNN Image Classifier** | TensorFlow · Keras · CNN · Transfer Learning · Data Augmentation

- Designed custom CNN architecture with 3 convolutional blocks achieving 92% accuracy on 10K+ image dataset; applied advanced data augmentation (flip, rotation, zoom, brightness) to reduce overfitting by 15% and tuned hyperparameters for 3× faster convergence.

**NLP Sentiment Analysis** | TensorFlow · LSTM · TF-IDF · Word2Vec · Bag-of-Words · N-grams

- Multi-class sentiment classifier processing 50K+ reviews at 88% accuracy; comparative benchmark of TF-IDF, Word2Vec, and BoW representations revealed TF-IDF + bigrams as the optimal strategy, yielding a 12% accuracy gain over baseline.

**Customer Churn Prediction & Segmentation** | XGBoost · K-Means · CI/CD · Docker · AWS EC2 ↗ GitHub

- XGBoost churn model (81% recall, 84% CV score) combined with K-Means clustering to surface 4 distinct customer segments; deployed via automated CI/CD on AWS EC2 with REST API for real-time scoring and business intelligence dashboard.

**Loan Approval Prediction System** | Logistic Regression · FastAPI · Docker · AWS EC2 ↗ GitHub

- End-to-end ML pipeline achieving 78.9% accuracy; FastAPI service deployed on AWS EC2 handling 500+ daily predictions with <100ms latency and containerized for reproducible scaling across environments.

## EDUCATION & CERTIFICATIONS

**Bachelor of Business Administration (BBA)** · Vivekanand Global University
*2024 – Present*

**IBM Deep Learning and Neural Networks with Keras** · Coursera
*2025*

**IBM Data Science Professional Certificate** · Python, SQL, Machine Learning, Data Analysis, Applied ML
*2024*